

Amanda: Unified Instrumentation Framework for Deep Neural Networks

Yue Guan*
bonboru@sjtu.edu.cn
Shanghai Jiao Tong University
Shanghai Qi Zhi Institute

Fan Yang
fanyang@microsoft.com
Microsoft Research

Yu Feng
yfeng28@ur.rochester.edu
University of Rochester

Yun Liang
ericlyun@pku.edu.cn
Peking University

Yuxian Qiu*
qiuyuxian@sjtu.edu.cn
Shanghai Jiao Tong University
Shanghai Qi Zhi Institute

Shuo Yu
y-shuo@sjtu.edu.cn
Shanghai Jiao Tong University

Yuhao Zhu
yzhu@rochester.edu
University of Rochester

Chen Zhang
chenzhang.sjtu@sjtu.edu.cn
Shanghai Jiao Tong University

Minyi Guo†
guo-my@cs.sjtu.edu.cn
Shanghai Jiao Tong University

Jingwen Leng†
leng-jw@cs.sjtu.edu.cn
Shanghai Jiao Tong University
Shanghai Qi Zhi Institute

Yunxin Liu
liuyunxin@air.tsinghua.edu.cn
Tsinghua University
Shanghai Artificial Intelligence
Laboratory

Lidong Zhou
lidongz@microsoft.com
Microsoft Research

Chao Li
lichao@cs.sjtu.edu.cn
Shanghai Jiao Tong University

Abstract

The success of deep neural networks (DNNs) has sparked efforts to analyze (e.g., tracing) and optimize (e.g., pruning) them. These tasks have specific requirements and ad-hoc implementations in current execution backends like TensorFlow/PyTorch, which require developers to manage fragmented interfaces and adapt their codes to diverse models. In this study, we propose a new framework called Amanda to streamline the development of these tasks. We formalize the implementation of these tasks as **neural network instrumentation**, which involves introducing instrumentation into the operator level of DNNs. This allows us to abstract DNN analysis and optimization tasks as instrumentation tools on various DNN models. We build Amanda with

two levels of APIs to achieve a unified, extensible, and efficient instrumentation design. The user-level API provides a unified operator-grained instrumentation API for different backends. Meanwhile, internally, we design a set of callback-centric APIs for managing and optimizing the execution of original and instrumentation codes in different backends. Through these design principles, the Amanda framework can accommodate a broad spectrum of use cases, such as tracing, profiling, pruning, and quantization, across different backends (e.g., TensorFlow/PyTorch) and execution modes (graph/eager mode). Moreover, our efficient execution management ensures that the performance overhead is typically kept within 5%.

Keywords: Deep Neural Network, Instrumentation

*Both authors contributed equally to this research.

†Jingwen Leng and Minyi Guo are corresponding authors of this paper.



This work is licensed under a Creative Commons Attribution-NonCommercial-NoDerivs International 4.0 License.

ASPLOS '24, April 27-May 1, 2024, La Jolla, CA, USA

© 2024 Copyright held by the owner/author(s).

ACM ISBN 979-8-4007-0372-0/24/04.

<https://doi.org/10.1145/3617232.3624864>

ACM Reference Format:

Yue Guan, Yuxian Qiu, Jingwen Leng, Fan Yang, Shuo Yu, Yunxin Liu, Yu Feng, Yuhao Zhu, Lidong Zhou, Yun Liang, Chen Zhang, Chao Li, and Minyi Guo. 2024. Amanda: Unified Instrumentation Framework for Deep Neural Networks. In *29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 1 (ASPLOS '24), April 27-May 1, 2024, La Jolla, CA, USA*. ACM, New York, NY, USA, 18 pages. <https://doi.org/10.1145/3617232.3624864>