

Quantized Sparse Training: A Unified Trainable Framework for Joint Pruning and Quantization in DNNs

JUN-HYUNG PARK, [Korea University](#), Republic of Korea

KANG-MIN KIM, [The Catholic University of Korea](#), Republic of Korea

SANGKEUN LEE, [Korea University](#), Republic of Korea

Deep neural networks typically have extensive parameters and computational operations. Pruning and quantization techniques have been widely used to reduce the complexity of deep models. Both techniques can be jointly used for realizing significantly higher compression ratios. However, separate optimization processes and difficulties in choosing the hyperparameters limit the application of both the techniques simultaneously. In this study, we propose a novel compression framework, termed as quantized sparse training, that prunes and quantizes networks jointly in a unified training process. We integrate pruning and quantization into a gradient-based optimization process based on the straight-through estimator. Quantized sparse training enables us to simultaneously train, prune, and quantize a network from scratch. The empirical results validate the superiority of the proposed methodology over the recent state-of-the-art baselines with respect to both the model size and accuracy. Specifically, quantized sparse training achieves a 135 KB model size in the case of VGG16, without any accuracy degradation, which is 40% of the model size feasible based on the state-of-the-art pruning and quantization approach.

CCS Concepts: • **Computing methodologies** → **Neural networks**; *Computer vision*; • **Computer systems organization** → **Embedded software**;

Additional Key Words and Phrases: Deep learning, model compression, joint pruning and quantization, neural network

ACM Reference format:

Jun-Hyung Park, Kang-Min Kim, and SangKeun Lee. 2022. Quantized Sparse Training: A Unified Trainable Framework for Joint Pruning and Quantization in DNNs. *ACM Trans. Embedd. Comput. Syst.* 21, 5, Article 60 (October 2022), 22 pages.

<https://doi.org/10.1145/3524066>

This work was supported by the National Research Foundation of Korea (NRF) grant funded by the Korea government (MSIT) (Grant No. 2021R1A2C3010430), the National Research Foundation of Korea (NRF) grant funded by the Korea government (MSIT) (Grant No. 2020R1A4A1018309), and Institute of Information & communications Technology Planning & Evaluation (IITP) grant funded by the Korea government (MSIT) (Grant No. 2019-0-00079, Artificial Intelligence Graduate School Program ([Korea University](#))).

Authors' addresses: J.-H. Park and S. Lee (corresponding author), [Korea University](#), 145, Anam-ro, Seongbuk-gu, Seoul, Republic of Korea, 02841; emails: {irish07, yalphy}@korea.ac.kr; K.-M. Kim, [The Catholic University of Korea](#), 43, Jibong-ro, Bucheon-si, Gyeonggi-do, Republic of Korea, 14662; email: kangmin89@catholic.ac.kr.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](https://permissions.acm.org).

© 2022 Copyright held by the owner/author(s). Publication rights licensed to ACM.

1539-9087/2022/10-ART60 \$15.00

<https://doi.org/10.1145/3524066>

ACM Transactions on Embedded Computing Systems, Vol. 21, No. 5, Article 60. Publication date: October 2022.